

ControlPlane.ai — Round 2 Final

Admission Control for AI That Acts

Accenture Innovation Challenge 2026 · Round 2 · dense hybrid (R2S1–R2S4)

ETERNAL: Default = *UNSUPPORTED* · entitlement = set-membership (zero LLM) · exact $R \times S$ matrix · one graph · hard gate on actions · FNR as typed format · two-pending-actions (R1 Edit + R3 Escalate **held**, never “blocked”) · UNKNOWN never → SUPPORTED · bias async-only · refuse-to-claim (about us)

Sources: docs/ARCHITECTURE.md · docs/reference/NARRATIVE.md · docs/reference/QA.md · docs/ps.md · Stages 1–4 absorbed in this FINAL (locks deleted after merge).

Official deliverable (docs/ps.md)	Artifact
Detailed Business Proposal	This document
Working Prototype	controlplane/ · examples/ · tests/
Pitch Presentation	R2S5.md submission/ControlPlane_Round2_Pitch.pptx +

Stage Check (compact)

Stage	Locked into	Status
R2S1 prototype scope	§§1–4	PASS
R2S2 enterprise envelope	§5	PASS
R2S3 build/demo	§6	PASS
R2S4 buyers/value/rollout	§§7–11	PASS

#	Eternal invariant	Status
1	Default = UNSUPPORTED	PASS
2	Entitlement = <code>span.ac1</code> $\subseteq$ <code>principal.clearance</code> ; zero LLM	PASS
3	Exact $R \times S$ matrix; never redrawn; no route parameter	PASS
4	One graph: STEP → SPAN → CLAIM → ACTION	PASS
5	Hard gate on actions , not tokens	PASS
6	Dual-action: R1 Edit + R3 Escalate (held , never “blocked”)	PASS
7	UNKNOWN never → SUPPORTED	PASS
8	FNR as typed format; empty until earned	PASS
9	Bias = async route-level only (never live matrix cell)	PASS
10	Refuse-to-claim list (about us)	PASS

1. Problem

Enterprises moved from AI that **answers** to AI that **acts**. Failure changed category: bad paragraph → **executed transaction**.

“Approved. Refund of INR 1,84,000 issued under clause 7.2 of the vendor agreement.”

Clause 7.2 does not exist. Failure is *absence* of evidence, not conflict. Filters pass it. Confidence can read 0.94. Money moved Tuesday, found Friday. If ungated, **the company wrongly pays out INR 1,84,000** — the customer did not lose money.

This is not “hallucination” as the market frames it. It is an **authorisation** problem: an unproven claim authorized an irreversible action. Solution is not better text scoring — it is **admission control**.

The system didn't fail. It was never asked to prove anything.

Everyone watches the exit. Nobody records the entrance.

Three structural reasons existing approaches fail:

1. **Inspect output, not context contract.** Without recording what the model was *given*, verification is unfalsifiable opinion. With it, verification is set-membership: *which span proves this claim?*
2. **Score the response, not the action.** Groundedness 0.82 means the same on a draft and a wire transfer. One threshold cannot price both.
3. **Identity-blind.** The common enterprise incident is a **correct answer to the wrong person**. No output-only inspector carries caller identity into verification.

Approach	Why it fails here
Post-hoc observability (LangSmith, Arize, ...)	Traces <i>after</i> commit — audit trail, not interlock
LLM-as-judge / wrappers	Opinion, identity-blind, too slow for commit, cannot state own error rate
Static guardrails	Miss lexically clean fabrications and ACL failures
RAG groundedness	Averages (one wrong figure drowns); action-blind; retrieval ≠ permission
Confidence / logprob	Named failure is <i>confidently</i> wrong
Composite risk scores	Three owners/costs collapsed into one number that maps to no actuator

Failure modes this plane addresses: (1) categorical claim with **no provenance** authorizing R3; (2) claim bound to evidence the **caller may not read**; (3) **two pending actions, different blast radii** on one response; (4) no real-time ground truth → **invert burden of proof**.

Overlap (brief)	Treatment on one graph
Fabricated personal detail, no span	Unsupported + no-span PII/secret shape; stays UNSUPPORTED
Correct HR fact, ACL excludes caller	Entitlement violation — independent of semantic correctness
Absent policy clause authorizing refund	Unsupported categorical on R3 → Escalate (held)

Overlap (brief)	Treatment on one graph
Bias on decision-shaped outputs	Not a per-response verdict; async flip-rate + CI

Multi-turn = more STEPs on the **same session ledger** — not a separate product. Prior assistant text is never evidence by reappearance. Unproven turn-1 claim used to authorize R3 in turn N remains unproven. Tool results enter as SPANs, not automatic authorization. Compounding risk = same unproven claim reaching higher R — already covered by the frozen graph.

Enterprise question is not “how do we score AI responses?” It is: *What evidence is required before a claim may authorize a specific action, for this caller, on this route?*

Buyer question a sceptic can answer without our slide: *What consequential actions does this route perform today, what is the loss if one is wrong, and what fraction can sit behind an earned admission boundary?*

2. Solution

ControlPlane.ai is an **admission-control layer** (firewall / transaction validator class — not observability, not a second model). Deployed as thin context-assembly SDK hook + OpenAI-compatible reverse proxy. No weights, logits, or fine-tuning. Integration cost is real — **and is the moat**.

AI output → claim proof → entitlement → R×S → action admitted / edited / escalated / held

One primitive, three reads:

STEP ■■■produces■■■ SPAN ■■■binds■■■ CLAIM ■■■authorizes■■■ ACTION

Performance reads forward. Cost reads backward (exact dead compute). Responsibility reads labels. One structure — not three detectors.

Load-bearing primitives

1. **Provenance outside the model** at context assembly — keystone: `source_id · ACL · content_hash · offsets` + calling principal. Model has **no write path**. Binding computed by the plane; model-emitted citations are not evidence; no open-web rescue.

2. **Default = UNSUPPORTED**. Earn SUPPORTED via recomputation or binding. Unsupported ≠ low confidence — **unproven**. UNKNOWN never → SUPPORTED.

3. **Claim-type routing**. Numeric/date/ID → recompute. Textual → bind (entailment, not string match). Derived/multi-hop → recompute or UNKNOWN.

4. **Entitlement = set-membership** (CALLER → CLAIM → SPAN → SOURCE ACL; `span.acl ⊆ principal.clearance`). Zero LLM. Lane 1 always on. Semantically correct + unauthorized still fails. **Most differentiated mechanism** — competitors cannot replicate without carrying identity into verification.

5. **Exact R×S matrix per pending action**. Disposition = **worst claim for that action — never an average**. R = irreversibility × audience × data class × autonomy. Action Interlock = sole final decider; pure rule engine.

6. **Hard gate on actions, not tokens**. Text hold-back ~150–300 ms. Speculative verification OK; **speculative release forbidden**.

7. **Escalate → evidence packet** (claim, spans, verdict, diff). Surgical Edit only (strip or one constrained regen; re-gate; second fail → Escalate).

8. **Publish own per-route FNR as typed format** — empty until earned. Emptiness is the credibility play.

Ungrounded / parametric: when no provenance spans exist for a route, it is declared ungrounded by construction; such claims **cannot authorize any action** regardless of blast radius. *We don't claim to verify what we were never given. We claim what we were never given cannot authorize an action.*

Refuse-to-claim (about us, not competitors)

1. **“We eliminate hallucinations.”** False if shipped. Narrower claim: ungrounded claims cannot authorise actions, and we report what we miss.
2. **“Zero integration, drop it in.”** We hook context assembly — real work, exact reason the design works. Scope = one SDK hook + OpenAI-compatible proxy (days on a standard retrieval stack, not quarters) — never sold as drop-in. Integration cost is the moat.
3. **“Zero added latency.”** We never make the model feel slow; we make the action wait. Quote only **≤40 ms p50 / ≤200 ms p95** — never 40 as p95.
4. **“One accuracy number across failure modes.”** Hallucination, leakage, and bias have different mathematics, error costs, and owners.

Operating claim: ControlPlane does not promise “safe” or “truthful” AI. It makes an unproven or unauthorized claim unable to authorize an action beyond the route's admitted boundary — and publishes what the plane missed.

3. Matrix

Exact frozen 16-cell R×S — transcribed, never redrawn, **no route parameter**:

	Contradicted entitlement violation	Unsupported categorical	Unsupported hedged	Unknown
R3	Block	Escalate	Escalate	Escalate
R2	Block	Edit	Edit	Escalate
R1	Edit	Edit	Pass + annotate	Pass + annotate
R0	Pass + annotate	Pass + annotate	Pass	Pass

R tiers: R0 = internal draft · R1 = user-visible read-only · R2 = reversible write / external send · R3 = irreversible or regulated (payment, deletion, publication, regulated advice).

Actuators (exact): Block · Edit · Escalate · Pass / Pass + annotate. Forbidden inventions: STREAM, Kill Span, Hold & Re-verify, Redact & Flag, COMMIT BLOCKED as R3 unsupported-categorical label. Autonomy downgrade / circuit breaker = spoken controls, not inventable demo labels.

Fail stance ≠ matrix: Matrix answers “evidence arrived and is bad → actuator?” Fail stance answers “evidence did not arrive in time?” Floors: R0/R1 fail **open + annotate**; R2/R3 fail **closed or escalate**. Universal fail-open forbidden.

Centrepiece cells (simultaneous on one refund response):

Pending action	Tier	Finding	Cell	Actuator
Show text	R1	Unentitled span grounds a claim (C3)	R1 × entitlement	Edit
Issue refund	R3	Clause 7.2 has no span (C2)	R3 × unsupported-categorical	Escalate — held with packet

Never collapse into one response-level verdict. **Never say “blocked” about the refund.** Same unsupported+category claim yields R1→Edit on a draft and R3→Escalate on a payment — proof scales with consequence; the matrix is not renamed severity.

4. Prototype

Exactly **two** live routes. Third live decision-support / bias route **refused**. Bias = proposal measurement only. Scope tension (resolved): brief names decision-support and bias; architecture defines bias as distributional/async — both stay in the **enterprise envelope**, never contaminate the live core.

Route	Dominant R	Mechanism proved	Live?
A. Refund agent	R1 text + R3 refund	Two-pending-actions	YES
B. Knowledge assistant	R0/R1	Provenance + entitlement flip	YES
Decision-support / bias	—	Async counterfactual	NO — envelope only

Refund risk signature: absent clause · numeric/ID error · entitlement leak into customer text · irreversible payment on same response · no real-time ground truth. Latency: R0/R1 text ≤ 40 ms p50 / ≤ 200 ms p95 (never quote 40 as p95); action verify amortised in tool RTT; text hold-back ~150–300 ms. Speculative verification OK; speculative release forbidden.

Knowledge: correct-but-unauthorized disclosure; fabrication; ACL-excluded span. Demo stages ACL-excluded → R1 × entitlement = Edit (not no-span PII → Block). Plane enforces carried ACLs — does **not** repair IAM. Authorization path: CALLER → CLAIM → SPAN → SOURCE ACL — set-membership, **zero LLM**.

Why this pair: Refund proves one graph, two actuators, hard gate, surgical edit, packet. Knowledge proves graph ≠ RAG groundedness: same claim, different principal → different outcome (RAG averages; ControlPlane prices worst claim per pending action). Together: *is this claim proven?* and *is this caller entitled to that proof?* A third live bias route adds async measurement without strengthening the dual-action centrepiece.

Assumptions (load-bearing): synthetic enterprise-shaped corpora only — no real PII; every span carries source/ACL/hash/offsets; generator via OpenAI-compatible API only (no weights/logits); decision policy = pure rule engine; traffic directional tens of thousands/week (prototype = curated single-node — no throughput claim); FNR null without trustworthy ground truth; corpus poisoning outside truth guarantee (hash makes source auditable, not factually true); DPDPA/GDPR as config packs — not certification.

Boundary — in

Prove	How
Provenance outside model	Spans before claims; model cannot author
One typed ledger	STEP → SPAN → CLAIM → ACTION
Default UNSUPPORTED	Earn via bind/recompute; UNKNOWN never → SUPPORTED
Claim-type routing	Numeric recompute; textual entailment; derived → UNKNOWN
Deterministic entitlement	Set-membership; zero LLM
Exact 16-cell matrix	Per pending action; worst-claim weighting
Dual-action centrepiece	R1 Edit + R3 Escalate held

Prove	How
Hard action gate	Mock refund committed: false while Escalate
Surgical Edit + evidence packet	Strip/regen once; packet = claim+spans+verdict+diff
Empty FNR schema	Typed nulls only — no fabricated %
Principal flip	Change only caller → outcome flips
Per-claim surface	Verified / Uncertain / Blocked; refund action = Held/Escalate
Ungrounded / parametric	No provenance → cannot authorize any action
Live compute	Binding/entitlement/interlock visible (~20–80 ms)

Boundary — out

Third live bias route · per-response bias · production load as mechanism proof · real payments/PII · live IAM repair · LLM-as-judge primary · confidence/risk scores · open-web truth · generative rewrite · fabricated FNR · Lane-3 on critical path · triage SLA UI · weights/logits · collapsing dual-action · calling refund “blocked.”

Measured results (this build)

All numbers are produced by `make eval` and `make bench` and regenerate on every run — they are measured, never asserted. The capability ledger and per-route Wilson CIs live in the generated README PDF (`submission/ControlPlane_Round2_README.pdf`) and `evals/last_run.json`.

This build earns the FNR on a 168-case self-authored corpus. The committed run (`evals/last_run.json`) reports:

Ungrounded FNR 1.1% (95% Wilson CI 0.2%–5.8%; n=94 ungrounded claims; 1 published miss).

We publish the interval — the honest claim is the interval, not a fake-perfect 0% and not a tighter bound than the data supports.

Passable-action FPR 0.0% (95% Wilson upper bound 15.5%; n=21 passable claims).

Hard-negative hold rate 64% [50.7%, 75.7%]. 53/168 cases are hard negatives (31.5%, above the 20% floor). We over-flag — that is our named next milestone. Clean strata hold 0/13.

One published miss: `struct-miss-000` — a response citing “Clause 4.1 permits this refund” where the span only says clause 4.1 *covers shipping delays*. Structural symbol lookup matches → SUPPORTED → low-tier action slips through. We show the judge exactly which case we got wrong.

No production proof. Production FNR is unknown until shadow replay over live traffic (ARCHITECTURE §7, §12). The honest claim: “on this self-authored corpus we miss 1.1% of ungrounded claims (95% Wilson CI 0.2–5.8%; n=94); production is unknown.”

Refuse-to-claim: we do not eliminate hallucinations; we do not claim drop-in integration; we do not claim zero added latency; we do not claim one accuracy number; we do not fill FNR with fabricated percentages.

Shape of the published claim (ARCHITECTURE §7): we hold **X%** of ungrounded responses at **Y ms** p50 — and here is the **Z%** we don't, with a Wilson interval.

5. Enterprise Envelope

Same plane on every route. Runtime:

```
App/Agent → Context SDK Hook (SPAN: source·ACL·hash·offsets)
→ OpenAI-compatible proxy
→ Evidence Ledger STEP→SPAN→CLAIM→ACTION
→ Action Interlock exact  $R \times S = f(R, S)$  [no route param]
→ Pass / Pass+annotate / Edit / Escalate / Block
→ Text hold-back · Action commit boundary
```

Shared (not soft-configurable): Provenance Recorder · Ledger · typed claims · proof paths · verdict set · Entitlement Auditor · Interlock · exact matrix · fail-stance floors · hard gate · surgical edit · packet · claim-level surface · FNR schema · Lane 1 always on.

RoutePolicy one-liner (configures budget, not truth):

```
RoutePolicy { route_id · tenant · use_case · provenance_scope · action_grammar
· action_to_R_mapping (locked R3: payment|deletion|publication|regulated_advice)
· verification_profile · fail_stance_by_R (may NOT loosen R2/R3 to fail-open)
· enforcement_mode shadow|canary|enforce · error_budget · escalation_target
· sampling_policy · geography_overlay (additive only) · latency_budget ≤40/≤200 }
```

Routes may change actions, R-mapping (except locked classes), Lane-2 budget, shadow/enforce, sampling, additive geo. Routes may **not** change: UNSUPPORTED default · entitlement · matrix cells · hard gate location · UNKNOWN→SUPPORTED · composite scores · Stage 1 dual-action scope. Low-consequence traffic gets **less verification budget and a proportionate actuator — not weaker truth semantics.**

Route profile	Dominant R	Lane posture	Enforcement
Customer-support refund	R1+R3 simultaneous	Lane 1 text; Lane 1+2 + interlock for R3	Shadow/canary before enforce; demo still shows hard gate on mock refund
Internal knowledge	R0/R1	Lane 1 dominant; ACL always inline	Enforce after shadow evidence
Decision-support	R0/R1 memo + R2/R3 action	Lane 2 derived; Lane 3 async bias	Action gate when earned; bias always async — not Stage 1 live

Lane	Runs	Path
1	Span membership, ACL, arithmetic, typed interlocks, PII shapes	Inline; always on; no LLM
2	NLI/binding for flagged textual / high-R	Near-line; timeout → UNKNOWN
3	Semantic-entropy, bias replay, calibration, shadow audit	Async; never critical path

Governance: versioned DAG of (signal, threshold, action, latency_budget); decision time = pure rule engine. Hierarchy narrows only: Enterprise baseline → geo overlay (additive) → tenant → route → action class. Overlays cannot loosen frozen invariants. Global parse-time invariants include: default UNSUPPORTED · UNKNOWN→SUPPORTED forbidden · entitlement always on · Lane 1 cannot disable · locked R3 classes · matrix immutable · R2/R3 fail closed/escalate · hard gate on actions · no composite score · no LLM at decision · bias never a matrix verdict · Stage 1 scope unreopened.

Policy lifecycle: Draft (hashed) → static validation → shadow replay → canary → auto-rollback if override >3× baseline or error budget breached → named-principal approval → promote. Day-one = **shadow**. Production enforce is **earned**.

Audit: principal · evidence · claim verdict · matrix cell · actuator · policy_version · verifier_versions · latency · lane · route_id — reconstructible from ledger + policy + frozen matrix.

Feedback = calibration, not adaptation. Learn thresholds/error-budget candidates via shadow→canary. Never learn: confidence/logprobs · silence-as-safety · aggregate override as permission to loosen · poisoned sources · prior assistant text as evidence · LLM-as-judge opinions · online weight updates. Override = evaluation evidence, not real-time security rewrite.

FNR schema (null until earned):

```
route_id · policy_version · evaluation_window · strata_definitions
· sampled_count_per_stratum · false_negative_count · ground_truth_positive_count
· FNR_estimate · CI_lower · CI_upper · ground_truth_method
· measurement_status # null|insufficient_sample|prototype_corpus|production_measured|stale
· limitations
```

Method: stratified shadow audit — 100% of Block/Escalate/Edit + sample of Pass; ground truth = human / expensive multi-verifier — **never** LLM-as-judge. Claim shape when earned: “*On this route we catch <measured>% of ungrounded claims at 40 ms p50 — and here is the <measured>% we don’t.*”

Bias (brief requirement, frozen stance): async route-level counterfactual flip-rate + CI; flag when CI excludes zero. **Never** claim → bias verdict → matrix. Never dropped; never a Stage 1 live route.

Layer	Enterprise	Stage 1 live	Proposal-only
Routes	Support + knowledge + decision + agentic	Exactly two	Extra packs
Core	Full plane	Dual-action + flip + ledger + empty FNR	Multi-tenant HA
Bias / FNR values	Async + measured when earned	Absent / null placeholders	Full programs
Scale	Tens of thousands/week (directional)	Curated single-node traces	Load/HA

Prototype = proof the plane works. Envelope = proof the plane can be operated. Pitch = both. Envelope residual risks (architecture): (1) derived/multi-hop false SUPPORTED → recompute-or-UNKNOWN + stratified FNR; (2) poisoned sources / wrong ACLs → hash + entitlement-by-source detector — defend claim↔evidence link, not source truth; (3) over/under-flag + R mis-mapping → shadow earn-out, locked R3, interlock in executor. Stage 2 IS configuration/governance/lifecycle around the primitive. Stage 2 IS NOT a second mechanism, redrawn matrix, live bias product, or trained safety model.

6. Build & Demo

Corpora (essentials — no breadth)

Source	ACL	Purpose
AGR-VENDOR-v3	agent-readable	Clauses 1–6 only; no clause 7.2 anywhere

Source	ACL	Purpose
ORD-1023	refund_agent	amount=184000 INR — C1 SUPPORTED
FIN-INTERNAL-NOTE	internal_analyst (excludes agent_refund_7)	C3 entitlement → Edit
INJECT-NOTICE	untrusted	Cannot author provenance
HR-COMP-L6	hr_partner	Principal-flip span

Principals: agent_refund_7 · analyst_01 · hr_partner_01. Canonical fixture: *“Approved. Refund of INR 1,84,000 issued under clause 7.2 of the vendor agreement.”* + one FIN-INTERNAL sentence. Span contract: span_id · source_id · ACL · content_hash · offsets · text · step_id.

Claim	Finding	Feeds
C1 amount/order numeric	Binds ORD-1023 → SUPPORTED	Proof works
C2 “under clause 7.2 ...” categorical	No span → stays UNSUPPORTED (absence ≠ contradiction)	refund.execute → R3 Escalate
C3 grounded on FIN-INTERNAL	ACL excludes caller → entitlement	text.show → R1 Edit

All three born UNSUPPORTED. Optional polish only after dual-action+flip green: paraphrase entailment · parametric ungrounded · numeric near-miss — cut polish before cutting dual-action.

Components

Real: Recorder · Ledger · Claim Extractor (fixture) · Numeric Recomputer · Entitlement Auditor · Interlock · Surgical Editor · Packet · Hold-back · Policy loader · Principal switch · FNR renderer · Trace Console · harness. Thin mock: generator stub · refund executor (REFUND HELD / committed:false — never COMMIT BLOCKED) · NLI labels. No LLM at decision time. No confidence field in the decision path.

Demo order (≤8 min; UI ≥60% ledger)

Governing test: if removing the graph leaves the demo looking the same, scope failed. Open on held transaction — never a risk statement. Core crisis ≤90s. First spoken sentence about failure contains **claim**, not response.

Must show	Must NOT show
Ledger ≥60%; spans before claims; matrix cells before actuators	Composite risk/confidence scores
Action Gate cold-open with committed boolean	“Response blocked” / COMMIT BLOCKED for R3 unsupported-categorical
Per-claim Verified/Uncertain/Blocked; refund = Held/Escalate	LLM-as-judge pane · open-web lookup
Evidence packet on every Escalate; empty FNR	Interlock bypass override · third-route chrome · bias widget · chatbot-majority

1. **Cold open gate:** refund.execute {amount:184000, reason:"clause 7.2", order_id:"ORD-1023"} · R3 · **HELD — ESCALATE** · executed:false
2. Expand ledger — spans with source·ACL·hash **before** claim verdicts
3. C1 SUPPORTED · C2 no span UNSUPPORTED · C3 entitlement violation

4. Highlight cells **before** actuators: `text.show` → `R1×entitlement→Edit`; `refund.execute` → `R3×unsupported-categorical→Escalate`
5. Surgical Edit strips C3; refund stays held; packet opens for C2 (claim, candidate spans [], verdict, diff)
6. Executor proves not committed — company does **not** wrongly pay out in the gated demo
7. Empty FNR schema visible (null placeholders only)
8. **Principal flip (≤2 min)**: `analyst_01` → Edit; flip only to `hr_partner_01` → Pass. Zero LLM.

Optional third beat (cut first): same unsupported claim as R1 vs R3 → actuator changes solely because R changed — or numeric 14 vs span 11 → R1 Edit.

Voice: `authorise` · `admit` · `prove` · `bind` · `refuse` · `hold` · `escalate` · `gate` — not `monitor` · `detect` · `observe` · `watch` · `guard` · `trust score` · `risk score`.

Run

```
python3 -m venv .venv && source .venv/bin/activate
pip install -e ".[dev]"
python3 -m pytest tests/ -v
python3 examples/refund_trace_demo.py
python3 examples/knowledge_flip_demo.py
```

Command	Proves
<code>pytest tests/ -v</code>	Matrix, entitlement, UNSUPPORTED default, hard gate, held≠blocked, R2S1 criteria on fixtures
<code>refund_trace_demo.py</code>	Dual-action: Edit + Escalate held; never collapsed; never “blocked”
<code>knowledge_flip_demo.py</code>	Same span/claim; outcome flips with clearance only

7. Buyers

Target users

Beachhead = high-consequence routes (refund-class R3 + mixed-governance knowledge) — not “all enterprise AI safety.”

Role	Who	Buys / operates
Economic buyer	Ops / CS / CHRO / CRO / CFO / CISO / Risk	Action authorisation + per-route error budgets. Responds to: <i>unproven claim cannot authorise an action</i> . Not: <i>we detect 95% of hallucinations</i> .
Technical buyer	Platform / ML infra / Identity	SDK hook + proxy + adapters; wants graph, not second black box; ≤40/≤200 ms
App / agent team	Route owners	Wire adapters; no app rewrite

Role	Who	Buys / operates
Day-to-day actor	Support / knowledge worker	Sees Verified/Uncertain/Blocked per claim; Held/Escalate on action — never enforces
Operator / governor	Route owners, SRE, Risk Ops	Shadow/canary/enforce; rollback; circuit breaker; ledger reconstruct
Escalation reviewer	Risk / ops	Evidence packet, not bare alert
Compliance / audit	Legal / DPO / audit	Influences via reconstructible trail

Split that matters: who **pays when it fails** ≠ who **runs it** ≠ who **types the answer**. Interlock enforces; plane publishes miss rate.

8. Impact Logic

Business case

No fabricated ROI. No “99%.” No net-savings slide. Value = **mechanism** → **consequence**; buyer fills middle terms. Tail risk, not average case.

Exposure = freq(consequential AI actions) × P(unproven/unauthorized claim) × loss/wrong action

Lever	Mechanism → consequence
A — Avoided wrong actions (primary)	R3 Escalate held (executed:false) is structural escape. Value = held true-positives × buyer’s direct cost of that action class; residual sized by published FNR. Artifact: attempted commit · matrix cell · gap · committed=false.
B — Blast-radius pricing	Lane 1 majority volume; expensive proof where R justifies. Same verdict annotates a draft and holds a payment. Answers brief’s one-size-fits-all latency problem.
C — Exact dead compute	Backward walk: STEP grounding zero accepted claims = waste (exact). Observability measures spend; this measures waste. Expose number; enterprise prices own traffic — no % saved on slide.
D — Alert fatigue without lowering gate	Matrix prices actuators; R0/R1 unsupported+hedged → Pass+annotate. Shadow earn-out. Metrics: overrides · gate-fail · edit/escalation per route .
E — Auditability	Hash-chained ledger + versioned policy. Reconstruct: action → cell → verdict → span → source/hash/ACL → principal → policy → latency. Regulator answer = pointer.
E2 — Unauthorized disclosure	ACL-excluded span caught when semantically correct. Does not fix IAM — stops IAM gaps being silently bypassed by a model .

Lever	Mechanism → consequence
F — Publish misses	Per-route FNR typed format; null until earned. Conversation: route · population · miss · CI — not “trust our safety score.”
G — Earned autonomy (secondary)	Shadow → enforce only where counterfactuals justify. More AI action after evidence — never lead claim.

9. Roadmap

Earn-out, not feature calendar. Enforcement earned, not switched on. Day-one = shadow. Misses published every phase.

Phase	What	Exit
0 — Prototype	Two live routes; synthetic corpora; ledger ≥60%; empty FNR	All R2S1/R2S3 criteria; dual-action crisis ≤90s
1 — Shadow production	1–2 real routes matching pattern; dual-emit; no production hold yet. Determinism works day one; statistics earn thresholds	Counterfactuals enough to open canary
2 — Canary R0/R1	Full policy lifecycle; enforce Edit/Pass+annotate on read-only	Tolerable ops; no material edit regression
3 — Limited R2/R3 enforce	Hard-gate selected actions; packets to humans; fill FNR only with trustworthy GT; locked R3 at parse	Buyer FNR/FP/override thresholds; committed=false on held TPs
4 — Broader envelope	More routes (decision-support template); additive geo; async bias; dead-compute FinOps; circuit breaker; fail stance tier-owned	Multi-route without second detector

10. Risks

Risk	Why it matters	Mitigation
False assurance on derived/multi-hop	Shallow entailment → false SUPPORTED — worse than no plane	Bypass NLI → recompute or UNKNOWN; UNKNOWN never SUPPORTED; timeout→matrix; FNR by claim type
Poisoned sources / wrong ACLs	Plane proves claim↔evidence + carried ACLs — not source truth / not IAM repair	Immutable source+hash; missing ACL→unentitled; entitlement-violation by source ; quarantine via policy
Prompt injection / model-declared bindings	Invent evidence edges	Plane computes bindings; model cannot author spans; disposition = rule engine

Risk	Why it matters	Mitigation
Over-flag → bypass; under-flag → liability	Classic guardrail death	Shadow default; blast-radius pricing; fail-stance floors; override>3× rollback; circuit breaker
R mis-mapping (payment as R1)	Wrong row without “breaking” matrix	Locked R3 at parse; interlock in executor , not only UI
Integration / provenance gaps	Incomplete spans → false UNSUPPORTED	Honest hook+proxy; evidence coverage measured; conservative high-R
Plane as SPOF	Outage / load-induced bypass	Tier-owned fail stance; universal fail-open forbidden
Feedback misuse / weight access	Train on overrides; require logits	API-only; feedback→policy candidates only; calibration ≠ trained judge
Switch-off after a quarter	Layers that interrupt low-blast text die	Gate on actions; majority Pass+annotate; earn-out before trust; integration cost stated as moat
Pattern-match as RAG/safety dashboard	Differentiation dismissed	Open on held INR 1,84,000; ledger majority UI; ban monitor/detect/risk-score vocabulary

11. Differentiation

ControlPlane	Everyone else
Provenance outside model at context assembly	Inspect output; treat model citations as evidence
Default = UNSUPPORTED	Default allow; flag what looks wrong
Entitlement = ACL set-membership	Identity-blind scorers
One graph, three reads	Three bolted tools
Exact R×S per pending action	Composite risk / confidence
Hard gate on actions	Gate on text/tokens
Publish per-route FNR (misses)	Publish precision (bother rate)

Ordered contrast

1. **vs observability** — post-hoc traces after harm. Observation without commit control is audit trail, not architecture. ControlPlane records entrance and interlocks commit; same graph yields exact dead compute.
2. **vs LLM-as-judge / static guardrails** — *does this look right? vs which span proves it?* Decision time = pure rule engine. Entitlement = identity, not classification. Default UNSUPPORTED is posture, not threshold tweak.
3. **vs RAG groundedness** — closest cousin, still short: retrieval-only, averages, action-blind. ControlPlane binds full provenance, prices **worst claim per action**, same unsupported claim → R1 Edit vs R3 Escalate. **Retrieval ≠ permission. Proof scales with consequence.**

Closer: Competitors publish precision. ControlPlane publishes misses and **refuses** to claim it eliminates hallucination/bias/privacy. Almost none disclaim themselves. Plane audited by the standard it enforces.

12. Fidelity + Content Laws

Invariant	Status
Default = UNSUPPORTED	Untouched
Entitlement / ACL set-membership; zero LLM	Untouched
Exact R×S; no route parameter	Untouched
Hard gate on actions; speculative release forbidden	Untouched
Dual-action R1 Edit + R3 Escalate held (never “blocked”)	Untouched
FNR typed format; empty until earned	Untouched
UNKNOWN never → SUPPORTED	Untouched
No LLM-as-judge / confidence as primary path	Untouched
Bias = async route-level only	Untouched
Refuse-to-claim (about us)	Untouched
Exactly two Stage 1 live routes	Untouched
Latency ≤40 ms p50 / ≤200 ms p95	Untouched
Surgical edit · evidence packet · Lane 1 always on · locked R3	Untouched

Law	Exact rule
Clause 7.2	Does not exist. Absence ≠ conflict. Never “caps/denies/doesn’t cover.” → Unsupported+categorical → Escalate , not Block.
Held ≠ blocked	Refund held and escalated with the evidence packet — never “blocked.”
Who pays	Company wrongly pays out. Customer did not lose money.
Dual action	Text → R1×entitlement→Edit (C3). Refund → R3×unsupported-categorical→Escalate (C2). Both simultaneous.
Latency	≤40 ms p50 / ≤200 ms p95. Never quote 40 as p95. Speculative verify OK; speculative release forbidden. Hold-back ~150–300 ms.
Refuse-to-claim	No: eliminate hallucinations · zero integration · zero added latency · one accuracy number.

Closing spine: AI can act → unproven claim must not authorize → provenance outside model → Default UNSUPPORTED → entitlement set-membership → R×S prices by consequence → hard gate on commit → publish misses.

Vocabulary: authorise · admit · prove · bind · refuse · hold · escalate · gate — not monitor · detect · observe · watch · guard · trust score · risk score · “responsible AI” as virtue.

Once you accept that an AI response is a set of claims requesting permission to act, you must capture provenance outside the model, invert the burden of proof, carry identity into verification, and gate the commit

path. Any softer design is a different product.

Appendix — Official brief coverage

Brief requirement (docs/ps.md)	Where	Freeze stance
Different risk/latency by use case	§5 RoutePolicy + profiles	Same graph/matrix; config only
Overlapping bias/hallucination/privacy	§§1, 5	Not one classifier
No reliable real-time ground truth	§§1–2; FNR null	Honesty over bluff
Over-/under-flag tradeoff	§§5, 8–10	Shadow earn-out; blast-radius pricing
Multi-turn / agents that act	§§1, 5 session ledger	No separate product
Evolving geo/industry regulation	§5 additive overlays	Cannot loosen matrix
API-only foundation models	§§2, 6	No weights/logits
Detection options (AI-as-judge, confidence)	§§2, 11	Reject as primary
Decision logic allow/edit/flag/block	§3 matrix	Transcribed
Architecture placement	§§2, 5–6	Pre-commit interlock + hold-back
Governance / audit	§5	Zero LLM at decision time
Feedback loops	§5	Calibration, not trained model
Metrics / FP/FN / trustworthiness	§§5, 8F	Publish misses as format
Multi use-case, tens of thousands/week, mixed governance	§§4–5	Directional; prototype simulated
Deliverables: proposal + prototype + pitch	Header map	Proposal=this Prototype=controlplane/ Pitch=next

End of ControlPlane.ai Round 2 Dense Final Hybrid — Stage Check PASS. Pitch from this file + live demo.