

ControlPlane.ai — Admission-Control Layer for AI That Acts

Accenture Innovation Challenge 2026 · Round 2 · Team ControlPlane · PS #1

Public GitHub: <https://github.com/Srujan0798/controlplane-ai> · branch main · tag v0.2.0-round2

1. Portal uploads (five fields)

These are the only artifacts the submission portal needs. This PDF is field 3.

#	Portal field	Artifact
1	Public GitHub link	https://github.com/Srujan0798/controlplane-ai
2	Prototype video	submission/ControlPlane_Round2_Prototype.mp4
3	README document (PDF)	submission/ControlPlane_Round2_README.pdf
4	Business proposal (PDF)	submission/ControlPlane_Round2_Proposal.pdf
5	Pitch deck (PPTX)	submission/ControlPlane_Round2_Pitch.pptx

2. What this is

Every AI response that asks permission to **act** — issue a refund, send an email, publish a record — is a set of **claims**. ControlPlane captures provenance **outside** the model (STEP → SPAN → CLAIM → ACTION) and admits or holds each action through a frozen interlock matrix. Unproven claims cannot authorize irreversible actions. Lane 1 is deterministic; no LLM sits on the critical path.

This repository is the working prototype. The numbers in the evidence section are produced by `make eval` and `make bench` and are regenerated on every run — they are measured, not asserted.

3. The mechanism

Run any response text through extract → bind → entitle → interlock → shadow:

- **extract** — Lane-1 segmenter produces typed claims (numeric / structural / temporal / textual / derived) with hedging and role-in-action, no LLM.
- **bind** — each claim is matched to provenance spans: BM25 + lexical gate (textual), Indian/Western numeric normalization (numeric), symbol-table lookup (structural), recompute (derived). Middle band → UNKNOWN, never a default pass.
- **entitle** — span ACL is checked against principal clearance; an unbound PII entity is a leak (Rule A) → Block at R2/R3.
- **interlock** — a frozen 16-cell matrix maps (blast-tier × verdict-column) to one of five actuators. The matrix is never redrawn.
- **shadow** — every decision is dual-emitted; the published FNR carries a Wilson interval.

4. Run it in 60 seconds (clean clone)

[illegible]

Refund language is **held and escalated with the evidence packet** — never 'blocked'.

5. What to look at

- **/gate** — paste a response; watch claims, binding method + rationale, symbol table, matrix cell, actuator, evidence packet, dead compute, per-stage latency.
- **refund autorun** — show_text → Edit · issue_refund → Escalate (held).
- **examples/refund_trace_demo.py** — same dual-action with fixtures off; amount via numeric.
- **evals/cases/** — labelled cases incl. hard negatives; this is the corpus a judge can attack.
- **controlplane/interlock.py** MATRIX — the sixteen transcribed cells, sole decider.

6. Capability ledger

Every mechanism in the architecture, marked honestly: **implemented and tested** / **implemented, prototype-grade** / **designed, not built**. This preempts 'is any of this real?' by answering first.

Mechanism	Status	Where
Claim extraction (Lane 1, no LLM)	implemented and tested	controlplane/extract.py
Numeric normalization (IN/US grouping)	implemented and tested	controlplane/numeric.py
BM25 + lexical entailment gate	implemented and tested	controlplane/bm25.py
Symbol table / structural lookup	implemented and tested	controlplane/symbols.py
Binder v2 route dispatch	implemented and tested	controlplane/binder.py
PII Rule A (leak → Block)	implemented and tested	controlplane/pii.py
Entitlement / ACL audit	implemented and tested	controlplane/entitlement.py
Frozen 16-cell interlock matrix	implemented and tested	controlplane/interlock.py
Lane deadlines (async gather)	implemented and tested	controlplane/lanes.py
Hash-chained evidence ledger	implemented and tested	controlplane/ledger.py
Override HTTP → chained ledger	implemented and tested	controlplane/server/app.py
Analyze arbitrary text	implemented and tested	POST /v1/controlplane/analyze
OpenAI-compatible FastAPI gate	implemented and tested	controlplane/server/app.py
Dual-action refund (Edit + Escalate held)	implemented and tested	examples/refund_trace_demo.py
Eval harness + Wilson CIs	implemented and tested	evals/run.py
Load bench n=10000	implemented and tested	scripts/load_bench.py
published_fnr from last eval	implemented and tested	GET /v1/controlplane/metrics
Shadow-replay threshold gate	implemented and tested	controlplane/feedback.py
Canary auto-rollback	implemented and tested	controlplane/feedback.py
Jurisdiction packs (EU/IN/GDPR)	implemented and tested	policies/*.yaml
Multi-turn claim inheritance	implemented and tested	controlplane/session.py
Economist / dead-compute walk	implemented and tested	controlplane/economist.py
Bias counterfactual (Lane 3)	implemented, prototype-grade	controlplane/bias.py
GPU / trained NLI adapter	designed, not built	Lane-2 slot (stub)
Production FNR on live traffic	designed, not built	shadow earn-out
Production IAM repair	designed, not built	out of Round 2 scope

7. Evidence (measured on this run)

Eval corpus: **168** labelled cases (426 bindings), **self-authored**. Ungrounded FNR **1.1%** (95% CI 0.2%–5.8%; n=94 ungrounded). We HOLD **98.9%** of ungrounded responses. Passable-action false-hold (FPR) **0.0%** (95% CI 0.0%–15.5%). Fail-closed caution on hard-negatives **64.2%**. Abstention (UNKNOWN rate) **0.5%** of bindings.

Published miss: struct-miss-000 (see evals/last_run.json).

Per-route binding distribution:

Route	n	supported	unknown	contradicted	unsupported	abstain
numeric	119	105	0	14	0	0.0%
structural	211	178	0	0	33	0.0%
textual	71	50	2	0	19	2.8%

derived	12	2	0	10	0	0.0%
temporal	11	11	0	0	0	0.0%
none	2	0	0	2	0	0.0%

Load bench: n=10000 at concurrency sweep [1, 4, 16, 64]. Per-request compute p50=0.4332 ms, p95=0.5113 ms, p99=0.6129 ms, mean=0.4455 ms, max=8.4518 ms. Best throughput 2258.8 rps. (In-process compute, not API latency; 40ms is a Lane-1 target, never a measured p95.)

Refund language in this system is **held and escalated with the evidence packet** — never 'blocked'. Irreversible actions are not released without provenance.

7.1 Honest reading of these numbers

Measured on a **168-case self-authored corpus**. Ungrounded FNR **1.1%** (95% CI 0.2%–5.8%), passable FPR **0.0%** (95% upper bound 15.5%). Hard-negative hold rate is **64%** (64.2%, 95% Wilson CI 50.7%–75.7%; 53/168 = 31.5% hard negatives) — we over-flag, we measured it, and it is our named next milestone. No production traffic; production FNR is unknown. The corpus ships in the repo so you can attack it.

7.2 What we refuse to claim

This list is about **us**, not about anyone else. **We do not claim:** elimination of hallucinations · zero integration cost · zero added latency · one accuracy number · production-scale proof.

8. Repo map

controlplane/ — pipeline (extract, binder, symbols, bm25, numeric, pii, entitlement, interlock, lanes, feedback, bias, shadow, server)

evals/ — corpus (cases/), harness (run.py), Wilson CIs, last_run.json

policies/ — jurisdiction packs (eu / in / gdpr / default)

scripts/ — bench, proposal/readme PDF builders, preflight, verify

examples/ — refund_trace_demo, knowledge_flip_demo, multi_usecase_demo

submission/ — five portal files: mp4 · README.pdf · Proposal.pdf · Pitch.pptx

Public GitHub — <https://github.com/Srujan0798/controlplane-ai>

Content laws (ARCHITECTURE §10) are enforced in CI: `pytest tests/test_content_laws.py`.